

LTFF Grant Renewal Proposal

TransformerLens maintenance + model-adapter breadth, 12 months

Applicant:	Bram Chen-Nakagawa
Funder:	Long-Term Future Fund
Grant manager:	Devante Kilmartin
Amount requested:	\$92,000 over 12 months
Disbursement:	4 tranches of \$23,000 on quarterly milestone acceptance, via Wise
Renewed:	2026-03

Track record (prior period)

TransformerLens is the standard open-source library for mechanistic interpretability of GPT-style models: a HookedTransformer API with activation hooks at every residual-stream, attention-head, and MLP position, plus ActivationCache. MIT licensed. ~3,680 stars, ~635 forks, ~2.4k Discord members.

- Monthly releases on a reliable last-Thursday cadence.
- Model-adapter breadth across GPT-2, Pythia, LLaMA-2/3/4, Qwen-2.5/3, Qwen-2.5-VL, Mistral-Small-3, Gemma-2/3, Phi-3, GPT-NeoX.
- Every adapter ships a weight-conversion acceptance test (float32 logits to 1e-4 absolute).
- Triage within a 72h PR / 5-business-day issue first-response SLA.
- Coordinated responsible disclosure for security-adjacent reports.

Proposed work

- Adapters for newly released public model families (published weights + architecture only).
- Maintenance of the hook surface and ActivationCache; careful cross-cutting changes with peer review.
- The monthly release train, docs for the adapter area, and the acceptance-test infrastructure.
- Compute (RunPod on-demand) for checkpoints that exceed the local 24GB box.

Budget (12 months)

category	allocated
Maintainer engineering time (stipend)	\$76,000
Compute (RunPod on-demand acceptance tests)	\$4,000
Tooling (Claude Max, HelloBonsai, domain)	\$2,000
Community / conference (ICML, NeurIPS, MATS)	\$6,000
Contingency	\$4,000
TOTAL	\$92,000

Terms acknowledged

Grant funds used only for the stated scope; deviation needs prior approval; missed milestones risk clawback of undisbursed tranches. Milestone reports state real, unrounded metrics. The grant ledger is kept separate from the consulting practice.